

KUBERNETES

IMRAN TELI

257) Introduction

258) Minikube for K8s setup

259) Kops for K8s setup

260) K8s Objects & K8s Documentation

261) Kubeconfig

Kubeconfig file

Use Kubeconfig file to organize information about

1) Clusters

2) Users

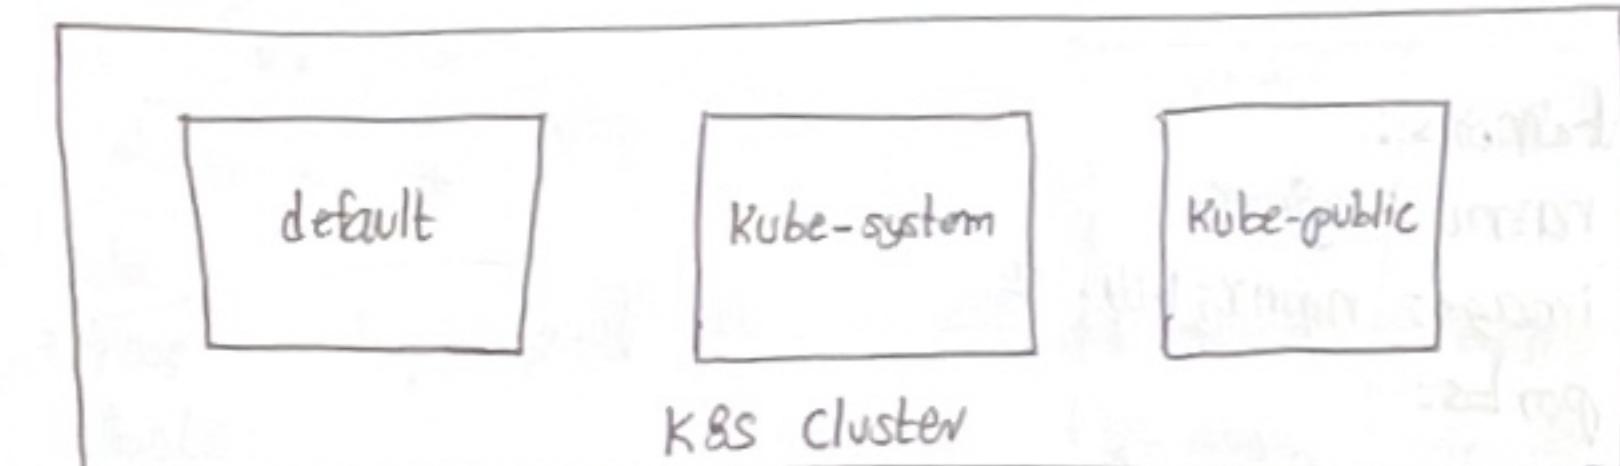
3) Namespaces

4) Authentication mechanisms

cat .kube/config

cat ~/.kube/config

262) Namespace



kubectl get ns (o) namespaces

We can isolate our environments using namespaces

=> Namespaces should be unique.

IMP Note:- We can gather resources in namespace. If we delete the namespace then our entire resources will delete with a single command.

```

# k get ns (or) namespace
# k get all (By default it will show all in default)
# k get all --all-namespaces
# k get svc -n mynamespace
# k create ns kubekart
# k run nginx --image=nginx -n kubekart

```

Create a Pod in particular namespace

```

# pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: kubekart

```

```

spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80

```

```
# k create -f pod.yml
```

```
# k get pod -n kubekart
```

If we delete namespace then entire things will delete

263) Pods:

A Pod is the basic execution unit of K8s application - the smallest and simplest unit in the K8s object model that you create or deploy. A Pod represents process running on your cluster.

Pods that run a single Container:

- ⇒ The one-container-per-pod model is the most common K8s usecase
- ⇒ Pod as a wrapper around a single container
- ⇒ K8s manages the pods rather than the containers directly.

Multi Container Pod

- ⇒ Tightly coupled and need to share resources
- ⇒ One main container and others as a sidecar or init container
- ⇒ Each Pod is meant to run a single instance of a given application
- ⇒ Should use multiple pods to scale horizontally.

pod.yml

```

apiVersion: v1
kind: Pod

```

metadata:

```

  name: webapp-pod
  labels:
    app: frontend
    project: infinity

```

Kind	Version
Pod	v1
Service	v1
Deployment	apps/v1
Ingress	networking/v1beta1

Spec:

```

  containers:
    - name: httpd-container
      image: httpd
      ports:
        - name: http-port
          containerPort: 80

```

GET & EDIT POD

```
# k get pod my-pod -o yaml
```

```
# k get pod my-pod -o yaml > my-pod.yaml
```

```
# k edit pod my-pod
```

264) Different Levels of Logging

```
# k get pods
```

Running → Fine

ImagepullBackoff → Image doesn't exist (or) Credentials issue

CrashloopBackoff → Pulling successfully, restart failed

↓
To get detailed description

```
# k get pod mypod -o yaml
```

```
# k logs podname
```

Delete Pod and create again

265) Service

Way to expose an application running on a set of pods as a network service. It is similar to Load Balancer.

There are three types of Services.

1) NodePort → Similar to port mapping in Docker. It is not for Prod. It is

2) Cluster IP only for exposing the application

3) ↳ If we don't want to expose externally. But internally
It is for internal communication.

3) Load Balancer: It is used to expose applications externally
for production environments

Please see Mumshad Mammambeth Diagram.

```
# service-def.yaml
```

apiVersion: v1

kind: Service

metadata:

name: webapp-service

spec:

type: NodePort

ports:

- targetPort: 80

port: 80

nodePort: 30005

protocol: TCP

selector:

app: frontend

```
# k create -f service-def.yaml
```

```
# k get svc
```

```
# k describe service servicename
```

```
# k get svc
```

Ports

8090: 30001/TCP

↓

↓

Internal Port External Port

To access application

→ Public IP of master → 3.139.133.219:30001

Cluster-IP

```
# kubectl get svc -n default
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
helloapp  ClusterIP  10.10.1.100  <none>       80/TCP    1d
```

LOAD BALANCER:-

```
# service-lb.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: helloworld-service
spec:
```

ports:

- port: 80
 targetPort: vproapp-port
 protocol: TCP

selector:

```
  app: vproapp
```

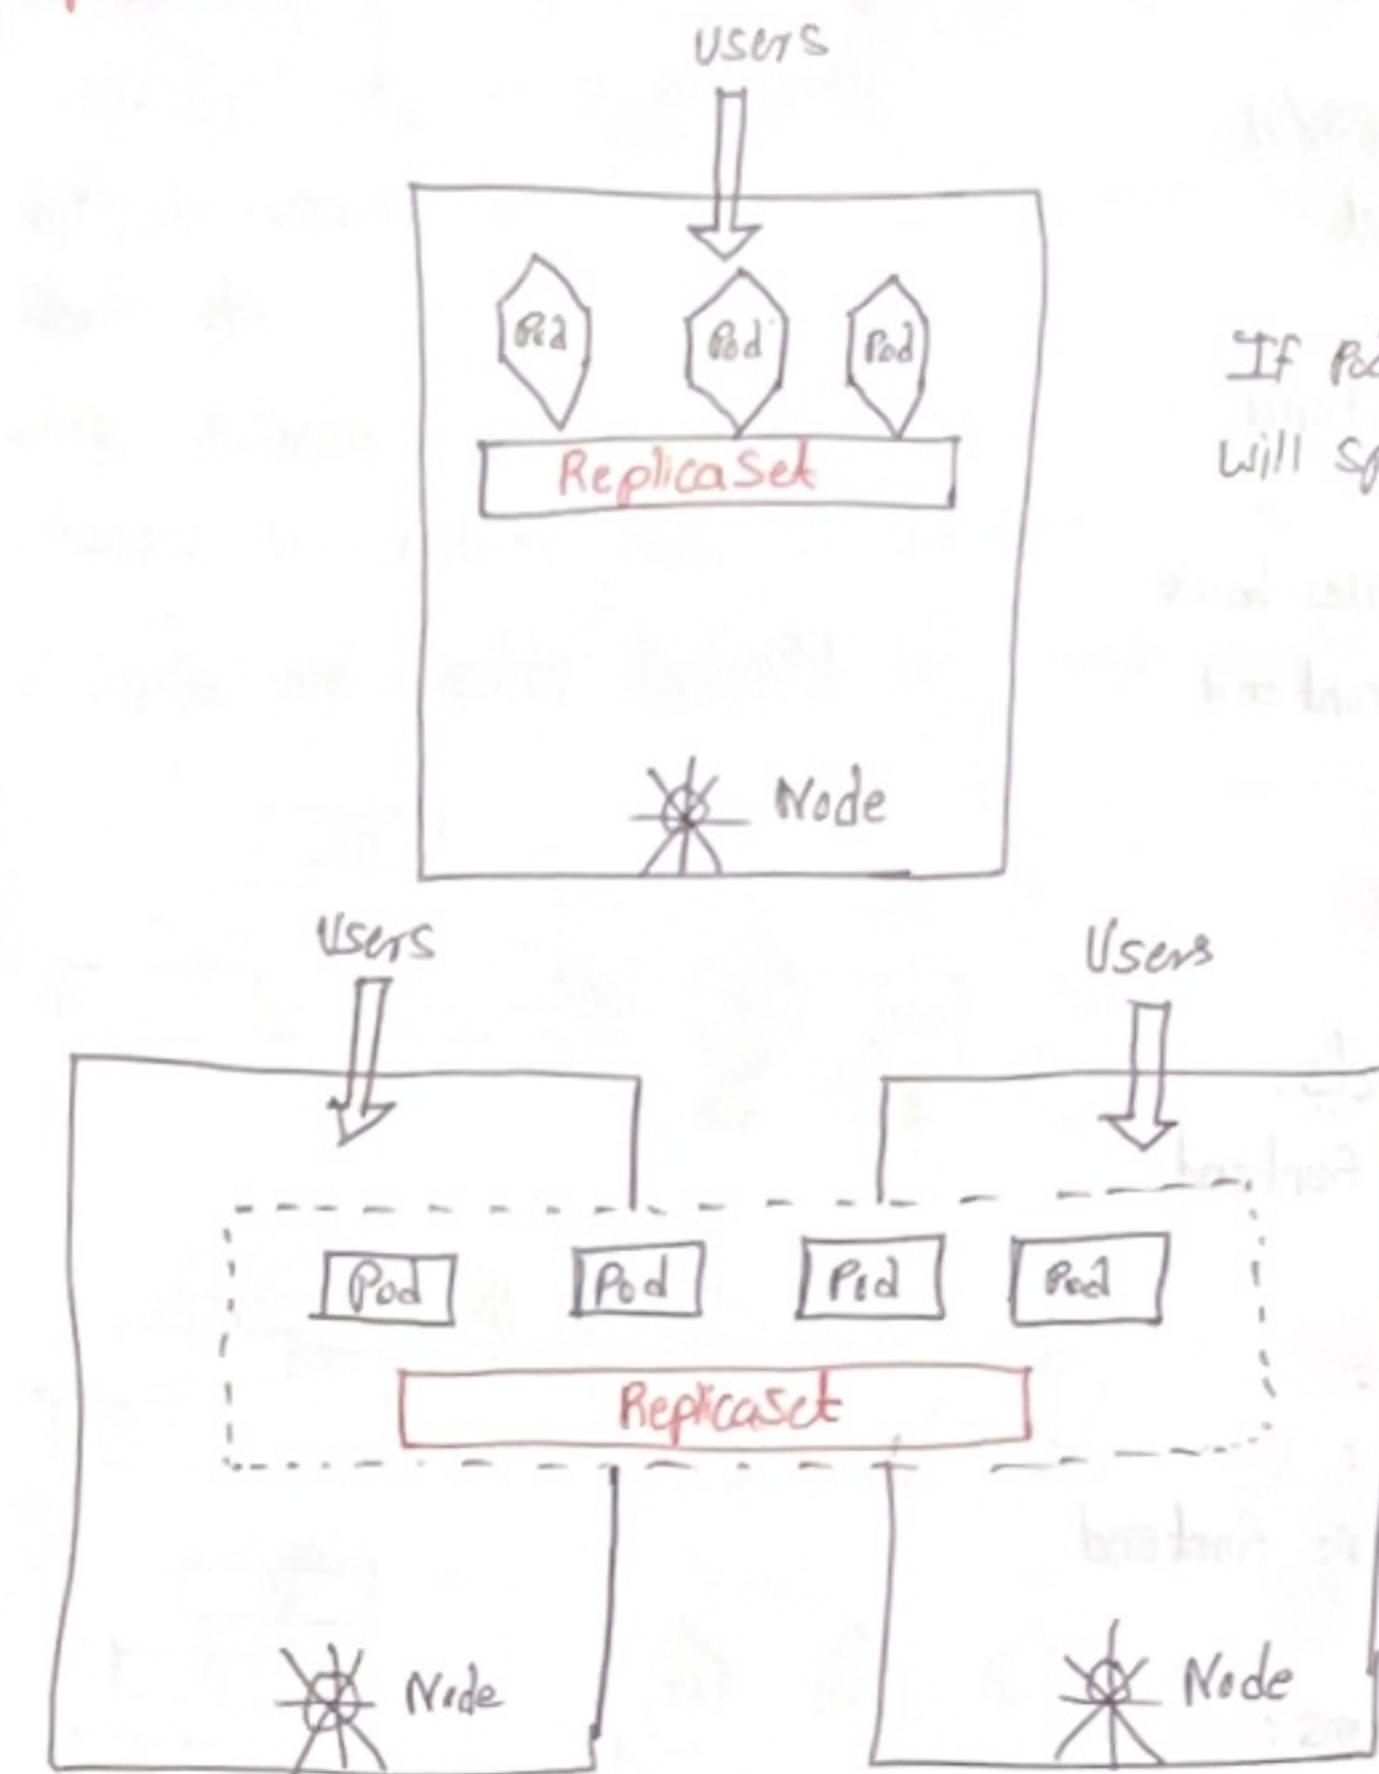
```
type: LoadBalancer
```

```
# kubectl get svc
```

```
# kubectl get svc
```

We will get LB IP in service External IP and the same is created in AWS console.

266) Replicaset



If Pod goes down the Replicaset will spin up the Pods

Replicaset

A Replicaset's purpose is to maintain a stable set of replica pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical pods.

Replicaset.yaml

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: frontend

labels:

app: guestbook

tier: frontend

spec:

replicas: 3

selector:

matchLabels:

tier: frontend

template:

metadata:

labels:

tier: frontend

spec:

containers:

- name: php-redis

image: gcr.io/google-samples/gb-frontend:v3

kubectl create -f replicaset.yaml

kubectl get rs

kubectl get pods

(2)

kubectl scale --replicas=1 rs/frontend

(2)

This is not recommended in
Production

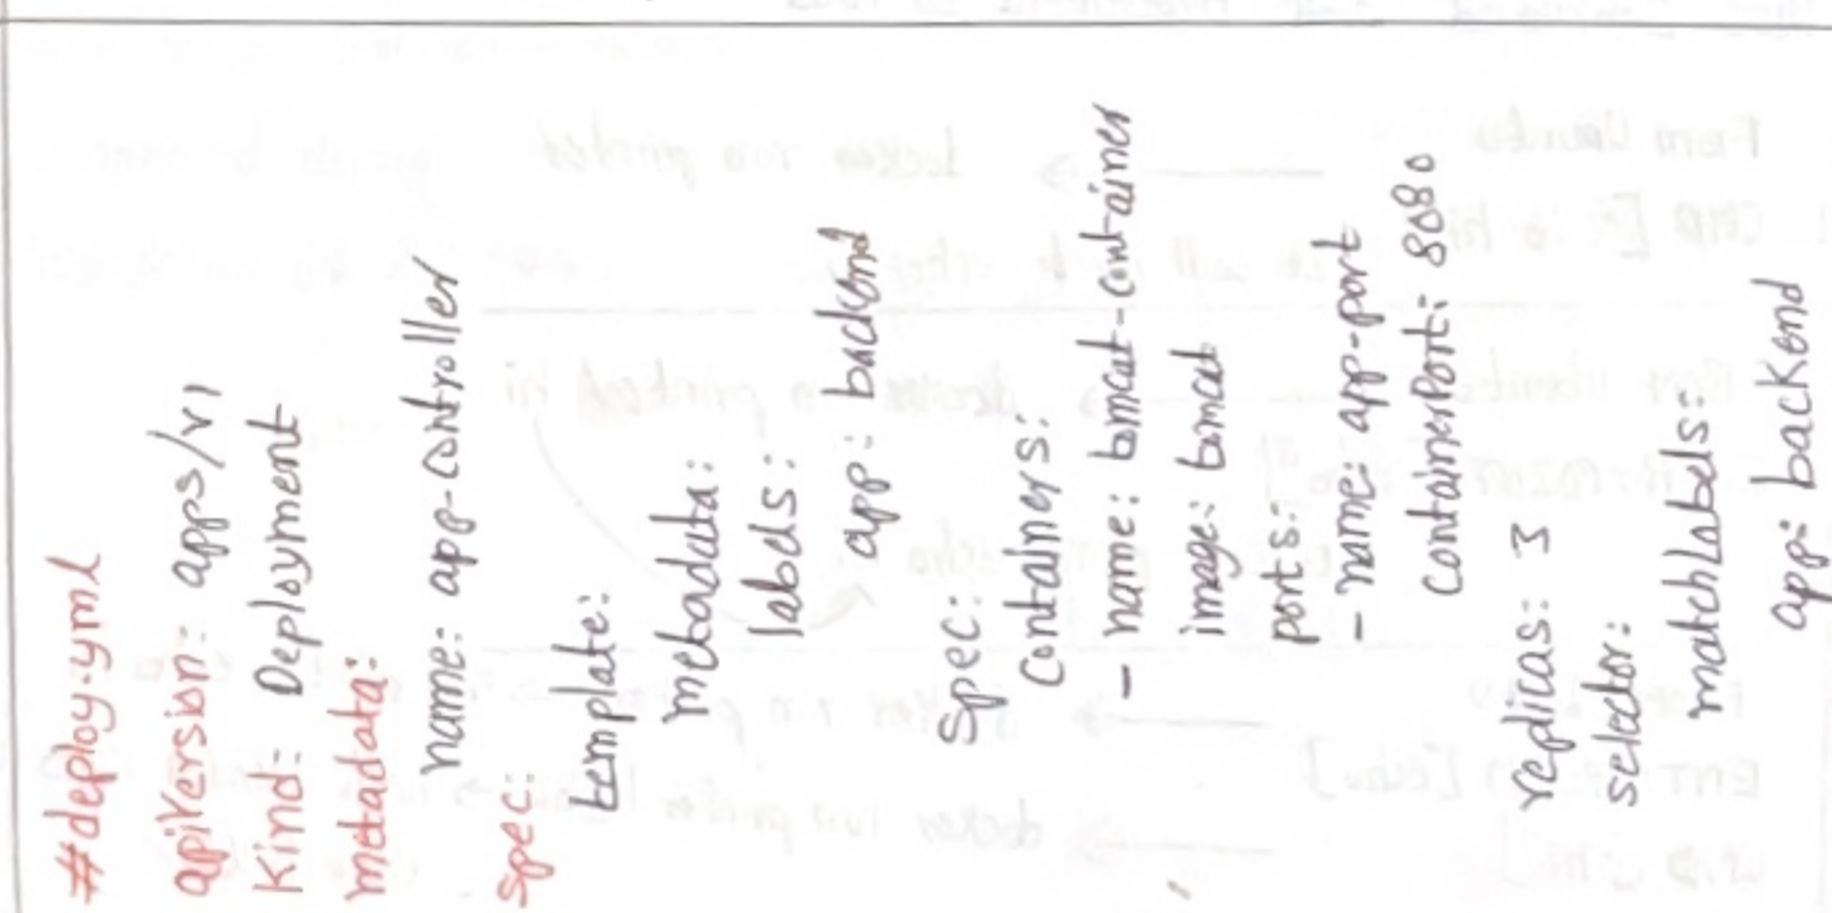
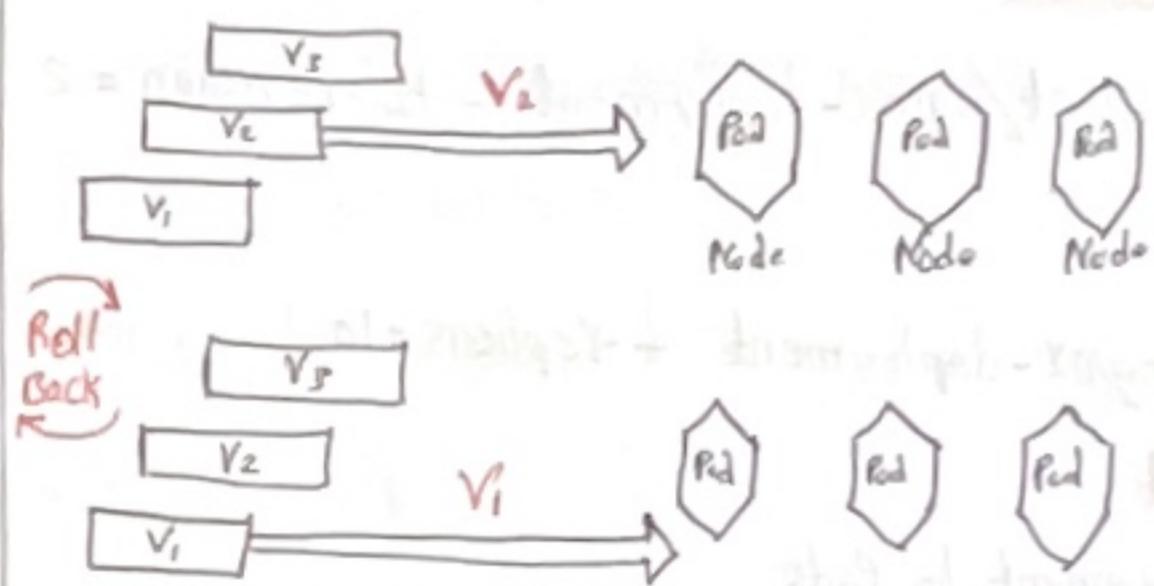
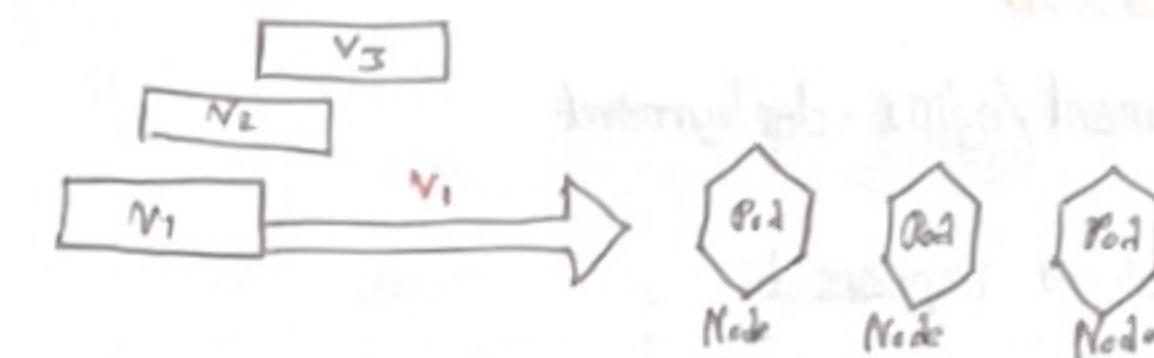
267) Deployment

Update, Rollback, changes gracefully

⇒ A Deployment controller provides declarative updates for Pods and ReplicaSets.

⇒ Define desired state in a Deployment and the Deployment controller changes the actual state to the desired state at a controlled rate

⇒ Deployment creates ReplicaSet to manage number of PODS



deploy.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: app-controlled

spec:

template:

metadata:

labels:

app: backend

spec:

containers:

- name: backend-container

image: backend

ports:

- name: app-port

containerPort: 8080

replicas: 3

selector:

matchLabels:

app: backend

269)

Volumes:

On disk files in a container are ephemeral, which presents some problems for non-trivial applications when running in containers. One problem is the loss of files when a container crashes. The kubelet restarts the container but with a clean state. A second problem occurs when sharing files between containers running together in a Pod. The k8s volume abstraction solves both of these problems. Familiarity with pod is suggested.

```
# Volume-pod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: dbpod
```

```
spec:
```

```
  containers:
```

- image: mysql:5.7

```
    name: mysql
```

```
  volumeMounts:
```

- mountPath: /var/lib/mysql

```
    name: dbvol
```

```
  volumes:
```

- name: dbvol

```
  hostPath:
```

```
    # directory location on host
```

```
    path: /data
```

```
    # this field is optional
```

```
    type: DirectoryOrCreate (If directory doesn't exist it will create)
```

```
# K apply -f volume-pod.yaml
```

```
# K describe pod dbpod
```

```
We have same error
```

Modified now

Delete Pod and create. Volume will mount

270)

Config Map

Environment Variables

Assign Variable values into a Pod

```
# pod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: db-pod
```

```
  namespace: dev
```

```
  labels:
```

```
    app: db
```

```
    project: infinity
```

```
spec:
```

```
  containers:
```

- name: mysql-container

```
    image: mysql:5.7
```

```
  env:
```

- name: MYSQL_DATABASE

```
    value: accounts
```

- name: MYSQL_ROOT_PASSWORD

```
    value: simopasswd
```

Config Maps:

Set & Inject variables / files in POD

Create Config Maps | Imperative

Create Config Maps | Imperative

```
# kubectl create configmap db-config --from-literal=MYSQL_DATABASE=accounts \
--from-literal=MYSQL_ROOT_PASSWORD=somecomplexpass
```

```
# kubectl get cm db-config
```

```
# kubectl get cm db-config -o yaml
```

```
apiVersion: v1
data:
  MYSQL_DATABASE: accounts
  MYSQL_ROOT_PASSWORD: passud
kind: ConfigMap
```

```
# kubectl describe cm db-config
```

Create Config Maps | Declarative

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: db-config
data:
  MYSQL_ROOT_PASSWORD: somepass
  MYSQL_DATABASE: accounts
```

```
# kubectl create -f db-cm.yaml
```

Pod Reading Config Maps :-

How to inject our config maps to the Pod's

apiVersion:

kind:

metadata:

spec:

containers:

- name:

image:

envFrom:

- configMapRef:

name: db-config

ports:

spec:

containers:

- name:

image:

env:

- name: DB-HOST

valueFrom:

configMapKeyRef:

name: db-config

key: DB-HOST

Example:- (See later)

27) Secrets

Share encoded / encrypted variables to Pod

Store and manage sensitive information, such as passwords.

Create Secrets | Imperative

```
# kubectl create secret generic db-secret --from-literal=MYSQL_ROOT_PASSWORD=  
=password
```

secret/db-secret created

Create files needed for rest of example

```
# echo -n 'admin' > ./username.txt
```

```
# echo -n 'f2d1e2e67df' > ./password.txt
```

```
# kubectl create secret generic db-user-pass --from-file=./username.txt  
--from-file=./password.txt
```

```
# kubectl get secret db-secret -o yaml
```

apiVersion: v1

data:

MYSQL_ROOT_PASSWORD: ab231gh2t1489ch==

→ Encoded

kind: Secret

metadata:

How to encode data:-

```
# echo -n "Secret pass" | base64
```

Decode:-

```
# echo 'd8bcd1245891208' | base64 --decode  
secretpass
```

It will be stored in k8s control plane. So it is safe

Create Secrets | Declarative:

```
# echo -n "Somepassword" | base64
```

```
C2afdzaklmYop==
```

```
# db-secret.yml
```

```
apiVersion: v1
```

```
Kind: Secret
```

```
metadata:
```

```
  name: mysecret
```

```
type: Opaque
```

```
data:
```

```
  my-root-pass: C2afdzaklmYop==
```

```
# k create -f db-secret.yml
```

Pod Reading Secret:

```
apiVersion:
```

```
Kind:
```

```
metadata:
```

```
spec:
```

```
  containers:
```

```
    - name:
```

```
      image:
```

```
        envFrom:
```

```
          - secretRef:
```

```
            name: mysecret
```

```
spec:
```

```
  containers:
```

```
    - name:
```

```
      image:
```

```
        env:
```

```
          - name: MYSQL_ROOT_PASSWORD
```

```
            valueFrom:
```

```
              secretKeyRef:
```

```
                name: mysecret
```

```
                key: my-root-pass
```

**

To pull image from your private registry

```
# echo -n "admin" | base64
```

```
YWRtaW4
```

```
# echo -n "mysecretpass" | base64
```

```
bXlzmnpad14789Nz
```

```
# secret.yml
```

```
apiVersion: v1
```

```
Kind: Secret
```

```
metadata:
```

```
  name: mysecret
```

```
data:
```

```
  username: YWRtaW4
```

```
  password: bXlzmnpad14789Nz
```

```
  type: Opaque
```

```
# k create -f secret.yml
```

Now Inject Into Pod, Pod.yml

```
apiVersion:
```

```
Kind: Pod
```

```
metadata:
```

```
spec:
```

```
  containers:
```

```
    - name:
```

```
      image:
```

```
        env:
```

```
          - name: SECRET_USERNAME
```

```
            valueFrom:
```

```
              secretKeyRef:
```

```
                name: mysecret
```

```
                key: Username
```

Continuation

```
- name: SECRET_PASSWORD
```

```
  valueFrom:
```

```
    secretKeyRef:
```

```
      name: mysecret
```

```
      key: password
```

```
      optional: false
```

```
  restartPolicy: Never
```

```
# k create -f pod.yaml
# k get pods
# k exec --stdin --tty secret-env-pod -- /bin/bash
    # echo $username
    # echo $SECRET_USERNAME
admin
    # echo $SECRET_PASSWORD
```

mysecretpass

Ingress

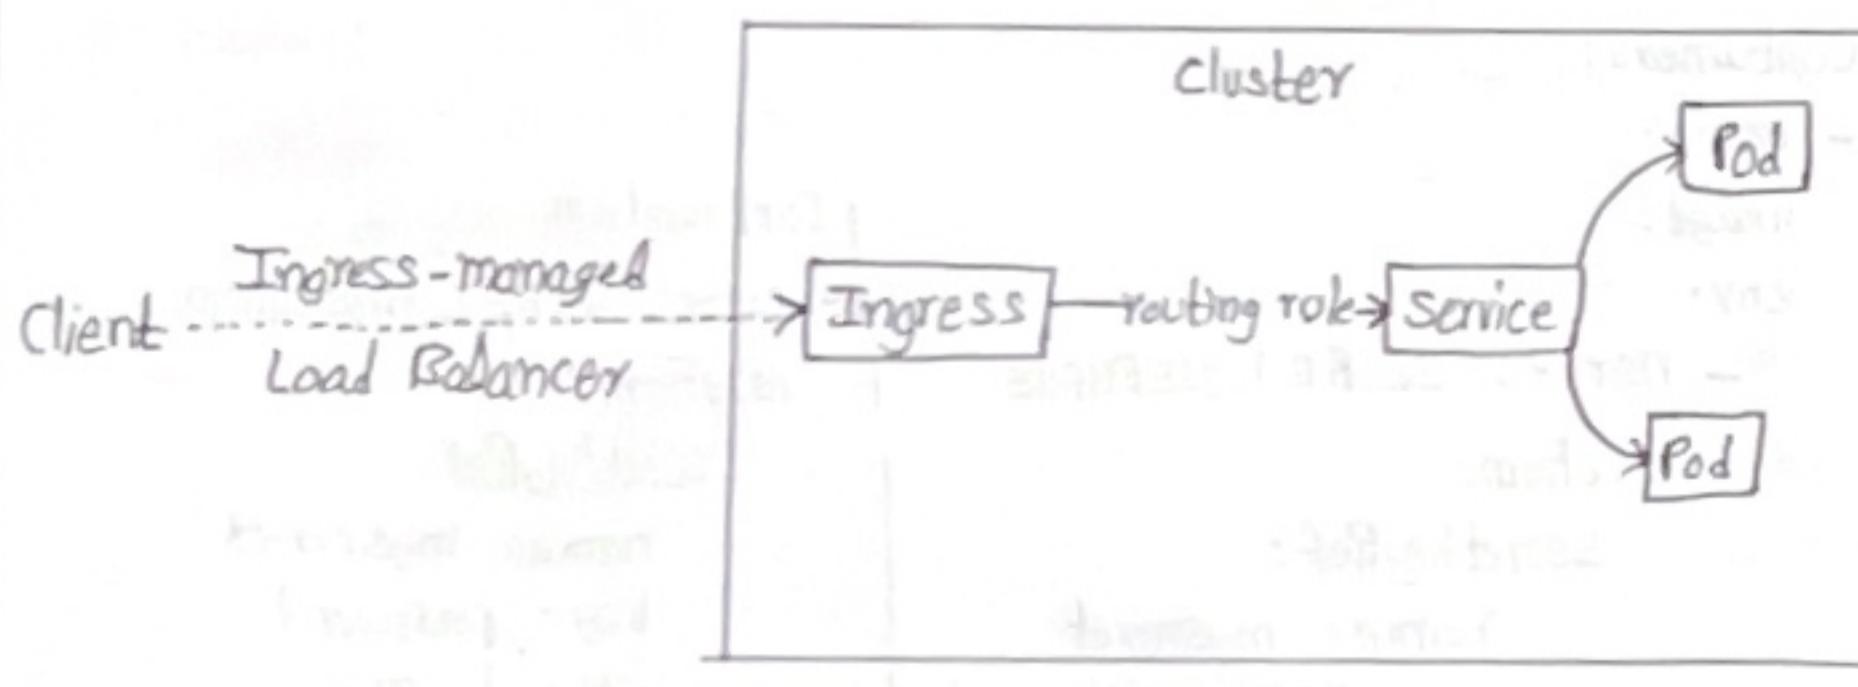
An API Object that manages external access to the service in a cluster, typically HTTP.

Ingress may provide load balancing, SSL termination and name-based virtual hosting.

What is Ingress?

Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster. Traffic routing is controlled by rules defined on the Ingress resource.

Here is a simple example where an Ingress sends all its traffic to one service.



Prerequisites:-

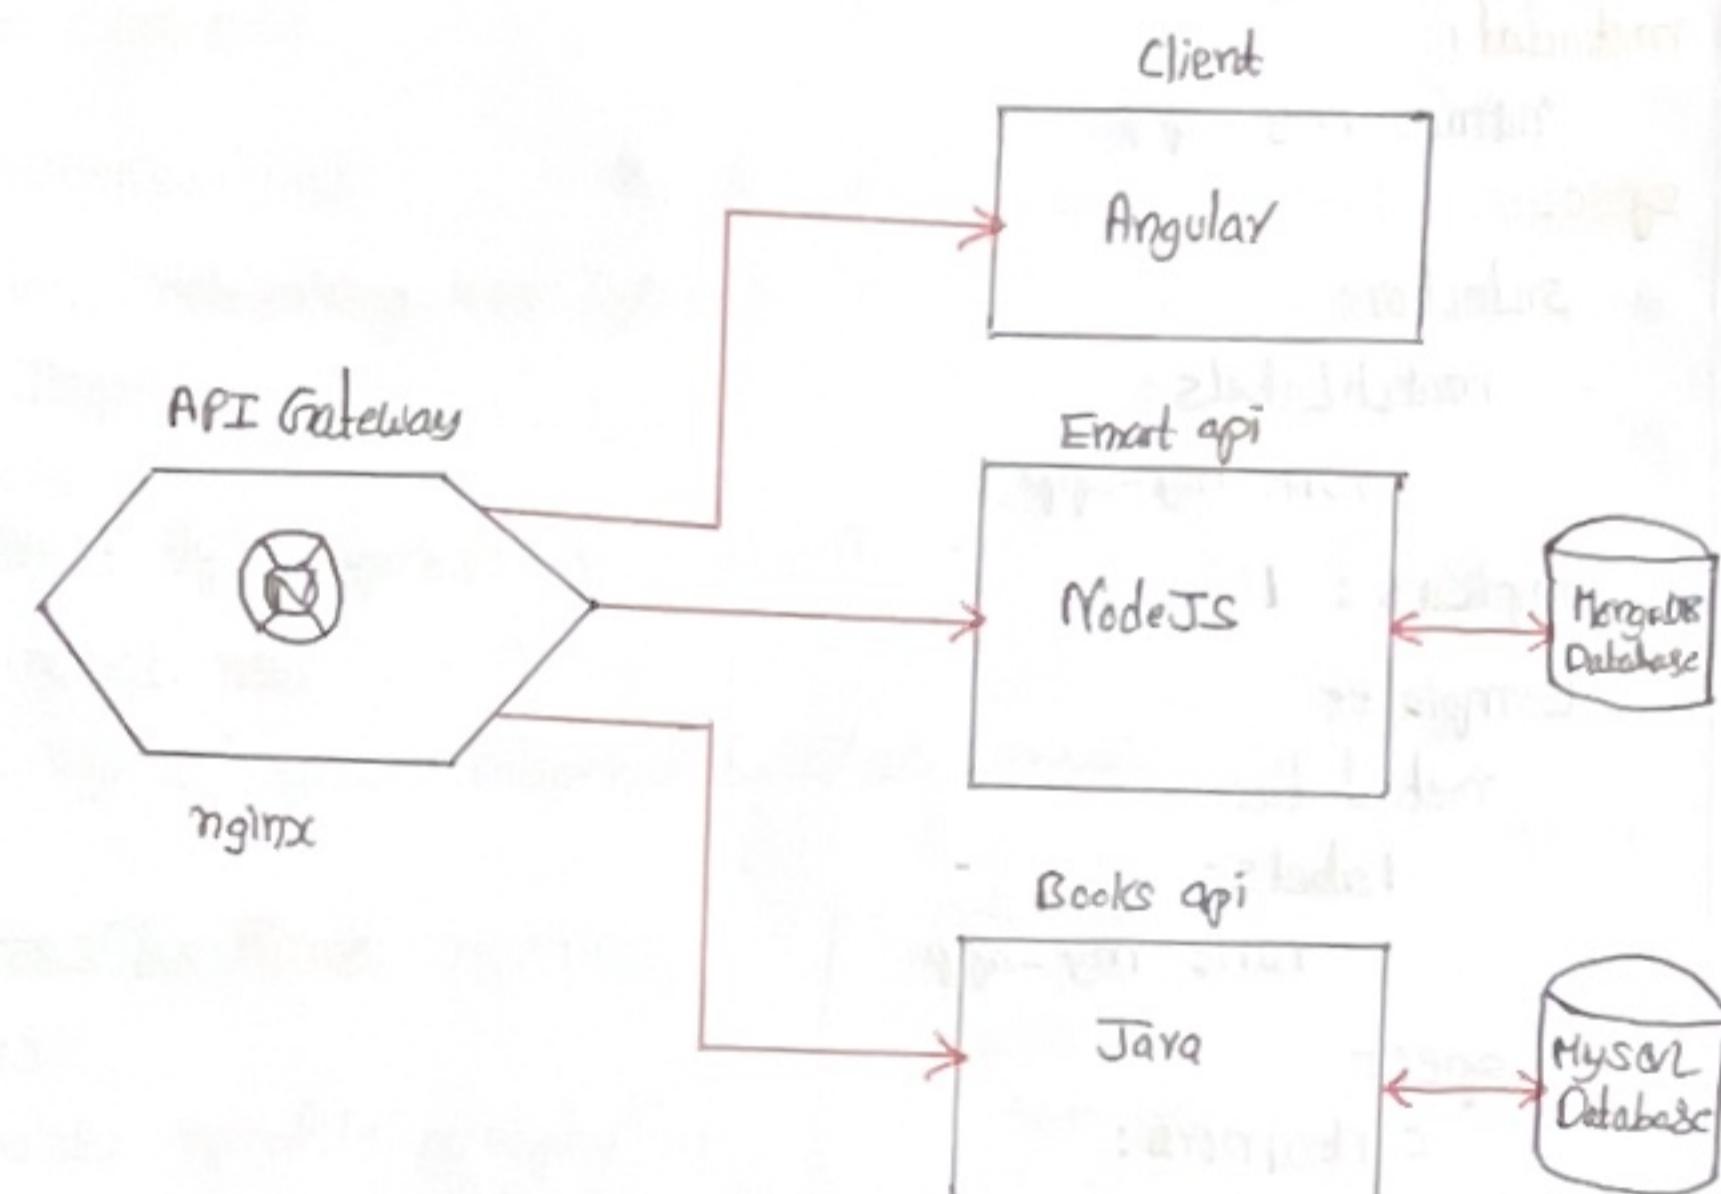
You must have an ingress controller to satisfy an Ingress. Only creating an Ingress resource has no effect.

You may need to deploy an Ingress controller such as Ingress-nginx. You can choose from a number of Ingress controllers.

Ideally, all Ingress controllers should fit the reference specification. In reality, the various Ingress controllers operate slightly differently.

We have different Ingress controllers

We will see NGINX Ingress Controller for Kubernetes



Install Ingress Controller:

Get Google Install nginx Ingress controller and run it

```
# K apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.3/deploy/static/provider/aws/deploy.yaml
```

```
# K get all -n ingress-nginx
```

Now we can see Load Balancer and user can access LB.

vprodep.yaml

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: my-app
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      app: my-app
```

```
  replicas: 1
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: my-app
```

```
  spec:
```

```
    containers:
```

```
      - name: my-app
```

```
        image: tomcat (or) our own image
```

```
    ports:
```

```
      - containerPort: 8080
```

vprocsvc.yaml

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: my-app
```

```
spec:
```

```
  ports:
```

```
    - port: 8080
```

```
      protocol: TCP
```

```
      targetPort: 8080
```

```
  selector:
```

```
    app: my-app
```

```
  type: ClusterIP
```

vpringress.yaml

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: vpro-ingress
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/use-regex: "true"
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: vprofile.groophy.in
```

```
      http:
```

```
        paths:
```

```
          - path: /login
```

```
            pathType: Prefix
```

```
            backend:
```

```
  pathType: Prefix
```

```
  backend:
```

```
    service:
```

```
      name: my-app
```

```
      port:
```

```
        number: 8080
```

```
# kubectl create -f vprodepl.yaml  
# kubectl create -f vprosvc.yaml
```

Now goto LB in AWS console and copy DNS and add DNS records with CNAME = value and add record.

```
# kubectl get svc  
# kubectl get describe svc my-app  
# kubectl apply -f vproxyingress.yaml
```

Just Overview

Controller

Deployment

Service

Created DNS Cname Record for LB

Created Ingress

kubectl get Ingress

Copy the URL and check in Browser (or) we can check in our Domain

Path Based Routing also available in docs please check video

273) Kubernetes CLI & Cheatsheet (Tricks)

Get k8s cheatsheet

274) Extras

Learn Taints & Tolerations from Hemeshad

Limits

Jobs [Kind: Job]

Cronjob

Daemonset

275) Lens

Lens can have central view of k8s cluster

Goto <https://k8slens.dev>

Install Lens